

Importing the needed libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px

from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, r2_score
from sklearn.preprocessing import LabelEncoder

from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score

import pickle
```

Loading The Dataset

```
In [2]: # Load the dataset into a pandas dataframe
dataset = pd.read_csv('Exam Results April.csv')
```

Analyzing The Dataset

```
In [3]: dataset.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 148 entries, 0 to 147
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype  
---  --
 0   Exam_Date   148 non-null     object  
 1   Section     148 non-null     object  
 2   Student_ID  148 non-null     int64   
 3   Class       144 non-null     object  
 4   Score       134 non-null     float64  
dtypes: float64(1), int64(1), object(2)
memory usage: 5.9+ kB

In [4]: dataset.shape
Out[4]: (148, 5)

In [5]: dataset.columns
Out[5]: Index(['Exam_Date', 'Section', 'Student_ID', 'Class', 'Score'], dtype='object')

In [6]: dataset.describe().T
Out[6]:
```

	count	mean	std	min	25%	50%	75%	max
Student_ID	148.0	107.450000	42.880014	1001.0	1037.75	1074.5	1111.25	1148.0
Score	134.0	85.201489	8.260587	70.0	78.00	84.5	92.00	100.0

Dropping Unwanted Data

```
In [7]: # Drop the "Exam_Date" column
dataset = dataset.drop("Exam_Date", axis=1)
```

Cleaning The Data

```
In [8]: dataset.isnull().sum()
Out[8]:
Section      0
Student_ID   0
Class        4
Score       14
dtype: int64

In [9]: sns.heatmap(dataset.isnull())
Out[9]: <AxesSubplot>
```

```
In [10]: dataset = dataset.dropna()

In [11]: dataset.isnull().sum()
Out[11]:
Section      0
Student_ID   0
Class        0
Score        0
dtype: int64

In [12]: sns.heatmap(dataset.isnull())
Out[12]: <AxesSubplot>
```

Finding The Average

```
In [13]: average_scores = dataset.groupby('Class')['Score'].mean()
average_scores

Out[13]:
Class
Arts      84.884615
Computer  86.520000
History   87.884615
Math      87.800000
Science   85.359256
Name: Score, dtype: float64
```

Finding Top 10 Student


```
In [14]: dataset.sort_values(by='Score', ascending=False, inplace=True)
top_students = dataset.head(10)

Out[14]:
```

	Section	Student_ID	Class	Score
20	B	1021	History	100.0
122	A	1123	Computer	100.0
118	B	1119	Arts	100.0
63	A	1054	Science	100.0
40	A	1041	Math	99.0
34	A	1035	Math	99.0
68	A	1069	Science	99.0
15	B	1016	History	99.0
145	B	1148	Computer	98.0
44	A	1045	Math	98.0

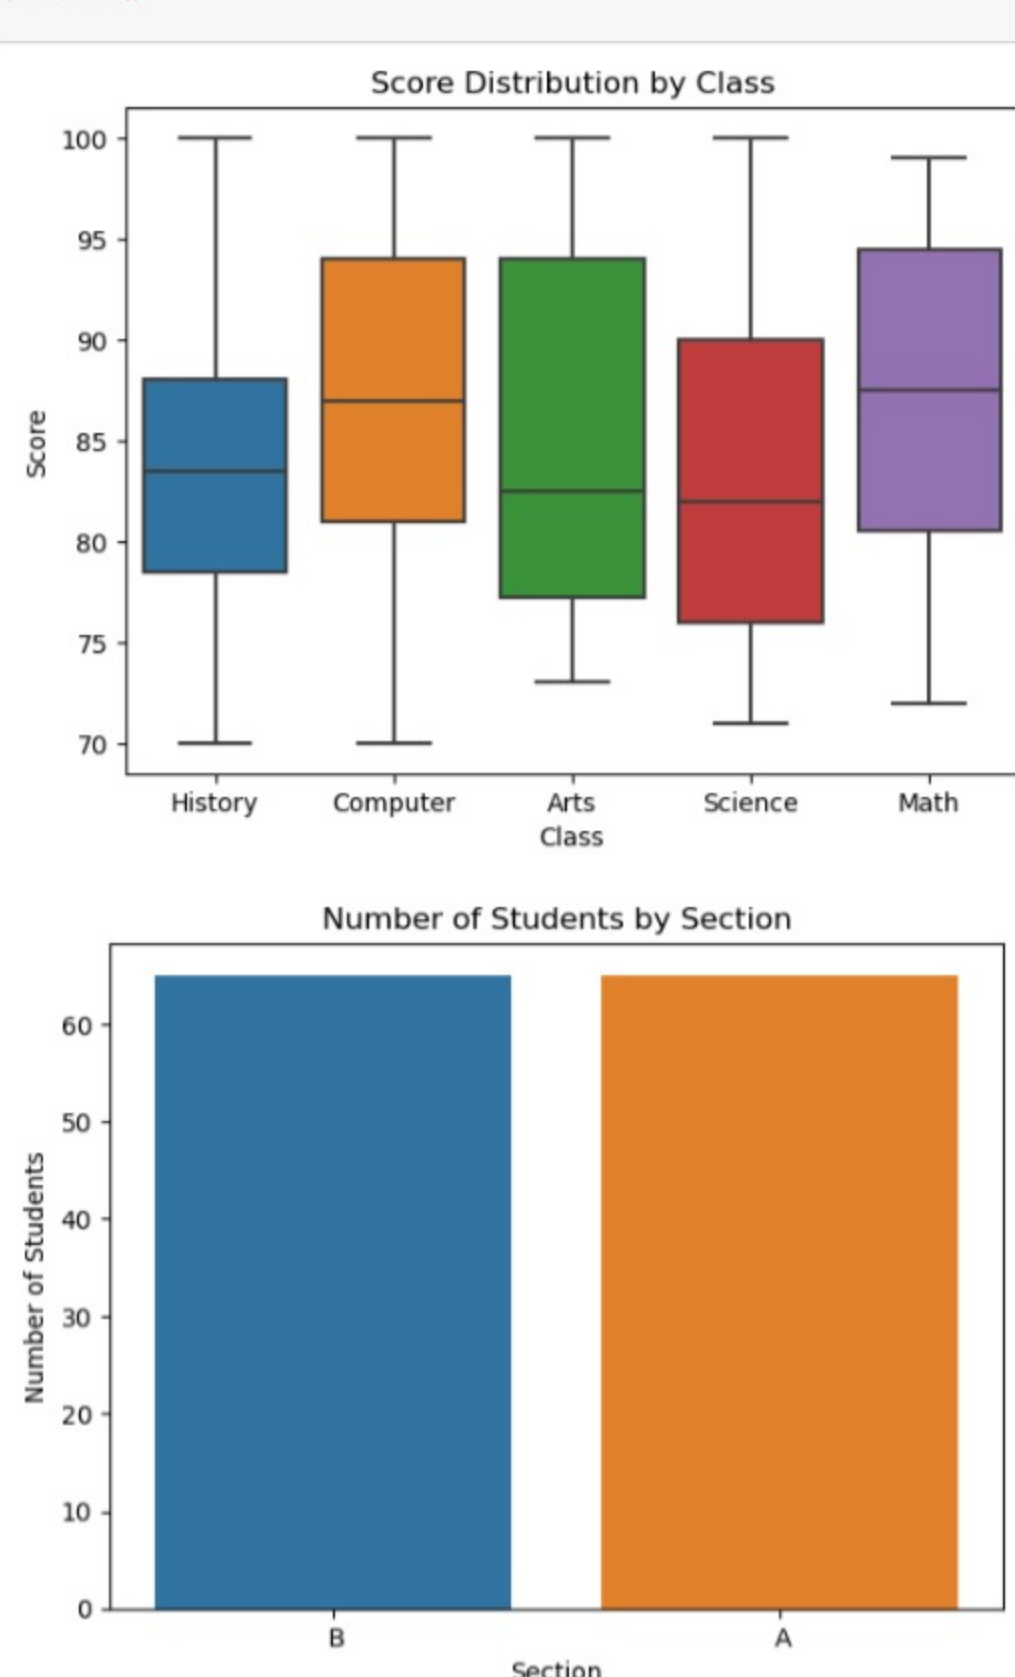
Visualizing the Data

```
In [15]: plt.bar(average_scores.index, average_scores.values, color=['black', 'red', 'green', 'blue', 'gray'])
plt.xlabel('Class')
plt.ylabel('Average Score')
plt.title('Average Scores by Class')
plt.show()
```



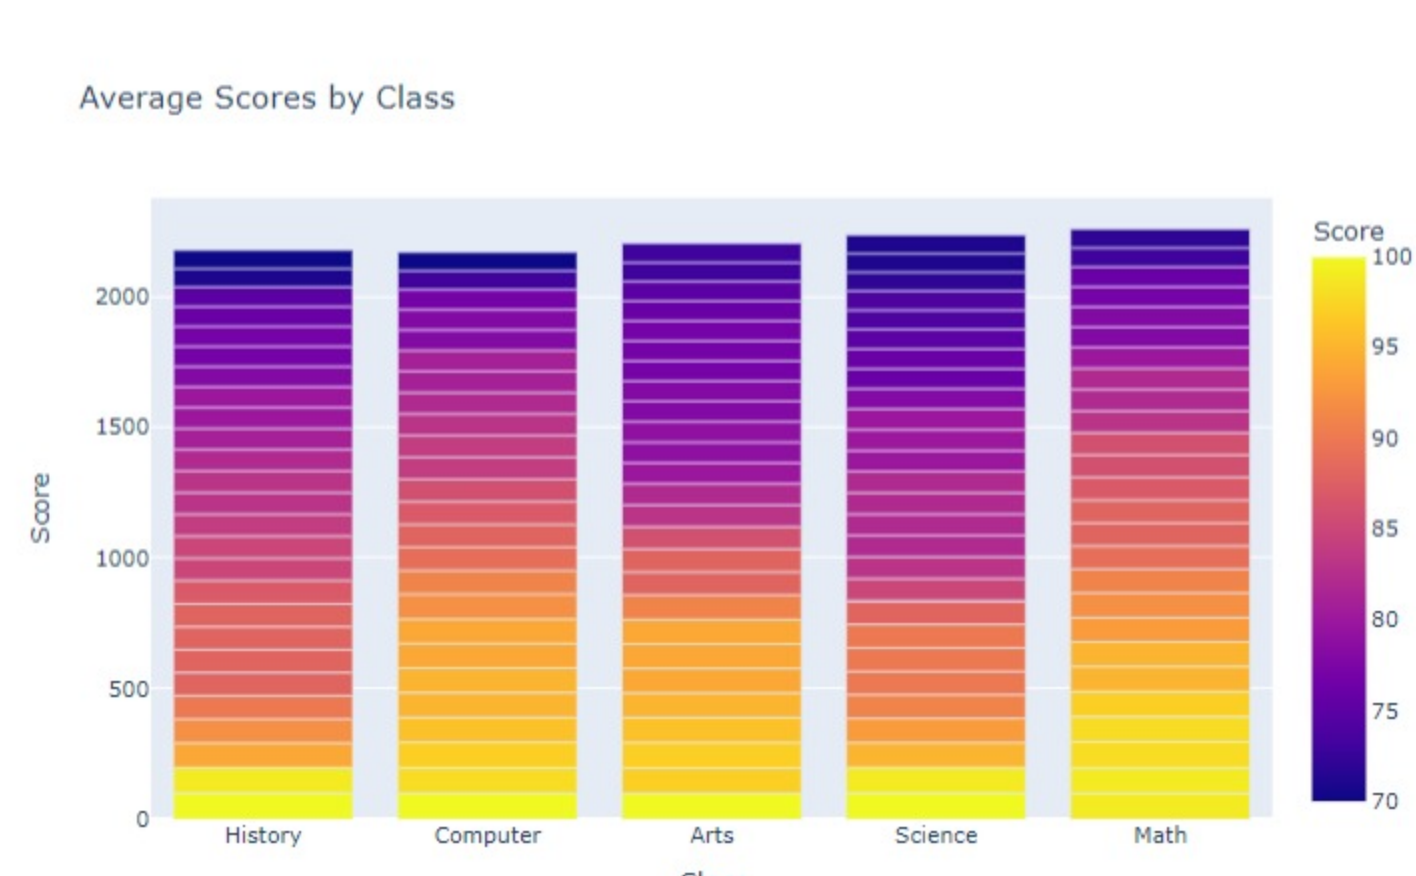
```
In [16]: # Create a bar plot showing the distribution of scores for each class
sns.boxplot(x='Class', y='Score', data=dataset)
plt.xlabel('Class')
plt.ylabel('Score')
plt.title('Score Distribution by Class')
plt.show()

# Create a bar plot showing the number of students in each section
sns.countplot(x='Section', data=dataset)
plt.xlabel('Section')
plt.ylabel('Number of Students')
plt.title('Number of Students by Section')
plt.show()
```



Using plotly.express

```
In [17]: fig = px.bar(dataset, x="Class", y="Score", title="Average Scores by Class", color="Score", width=800)
fig.show()
```



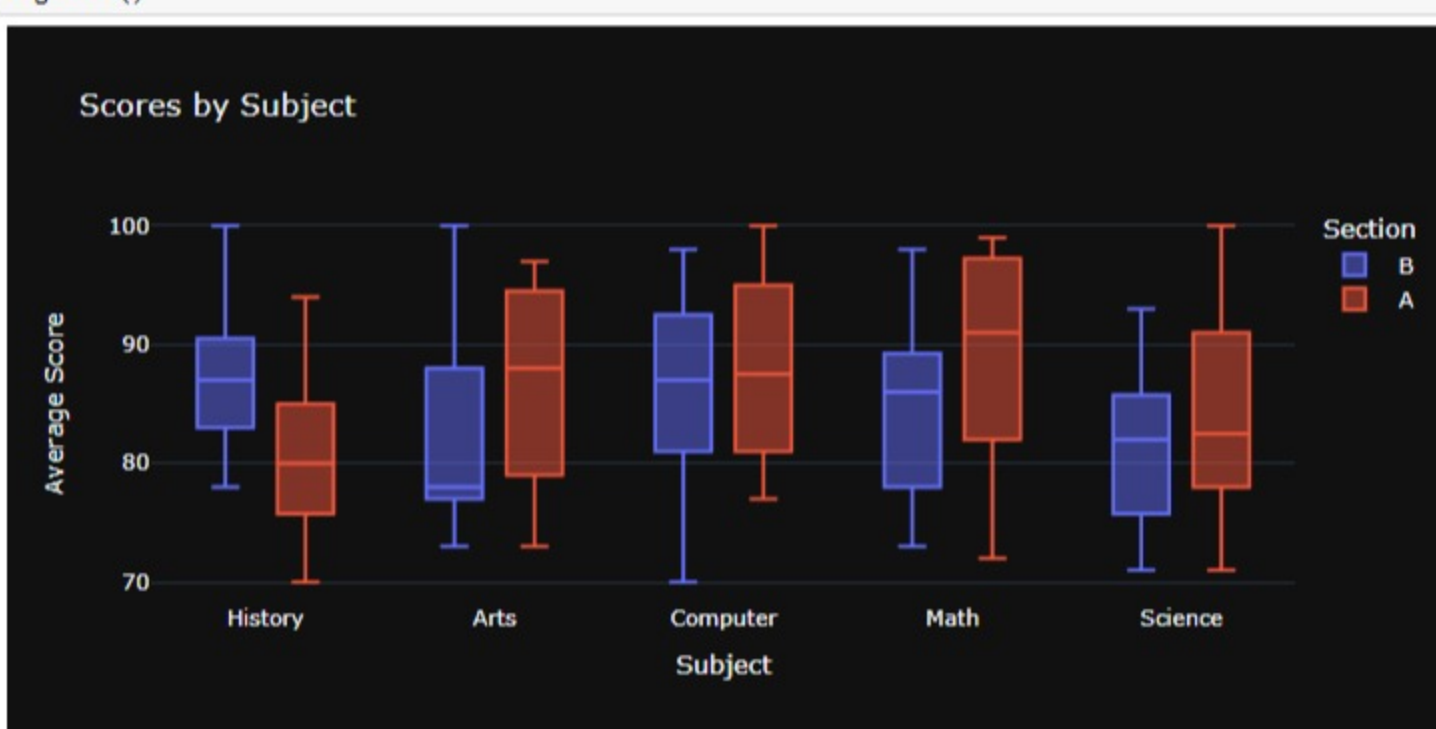
```
In [18]: fig = px.scatter(dataset, x="Score", y="Class", title="Scores by Subject", color="Section",
                        hover_name="Student_ID", labels={"Score": "Average Score", "Class": "Subject"},
                        template="plotly_dark", width=800, height=400)

# Show the plot
fig.show()
```



```
In [19]: # Plot the data using a box plot
fig = px.box(dataset, x="Class", y="Score", title="Scores by Subject", color="Section",
             hover_name="Student_ID", labels={"Class": "Subject", "Score": "Average Score"},
             template="plotly_dark", width=800, height=400)

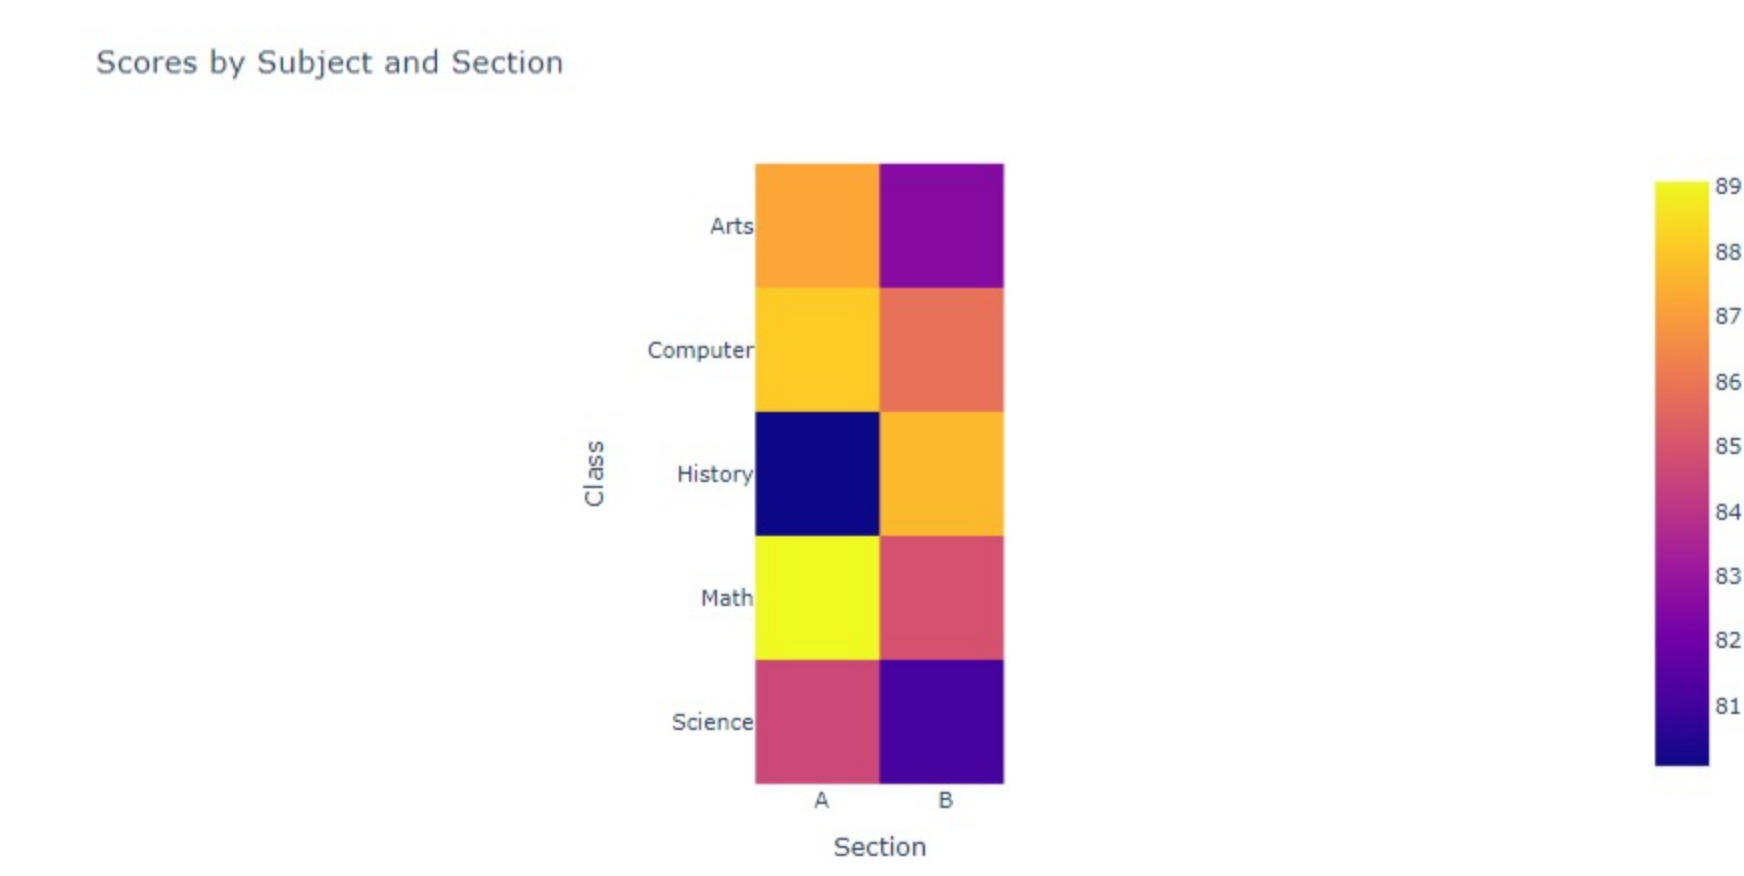
# Show the plot
fig.show()
```



```
In [20]: # Create a pivot table to summarize the data
pivot = dataset.pivot_table(index="Class", columns="Section", values="Score")

# Plot the data using a heatmap
fig = px.imshow(pivot, title="Scores by Subject and Section")

# Show the plot
fig.show()
```




Finding the Correlation

```
In [21]: # Calculate the correlation matrix
correlation = dataset.corr()

# Visualize the correlation matrix using a heatmap
sns.heatmap(correlation, annot=True, cmap="Pastel1_r", linewidths=0.5,
            fmt=".2f", cbar=False, square=True, xticklabels=correlation.columns, yticklabels=correlation.columns)

Out[21]: <AxesSubplot>
```



Machine Learning Over The Data

```
In [22]: # Encode the "Class" column to numerical values
le = LabelEncoder()
dataset["Class"] = le.fit_transform(dataset["Class"])
dataset["Section"] = le.fit_transform(dataset["Section"])

In [23]: # Select the features and target
X = dataset[["Class", "Section"]]
y = dataset["Score"]

In [24]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

In [25]: # Create a random forest regressor
model = RandomForestRegressor(n_estimators=100, random_state=42)

In [26]: model.fit(X_train, y_train)

Out[26]: RandomForestRegressor(random_state=42)

Save the Model

In [28]: # Save the model to a file
with open("model.pkl", "wb") as f:
    pickle.dump(model, f)

Load the Model

In [29]: # Load the saved model
with open("model.pkl", "rb") as f:
    saved_model = pickle.load(f)

Test the Model

In [30]: # Make predictions on the test data using the saved model
y_pred = saved_model.predict(X_test)

Out[30]: array([[86.1228307, 80.880349802, 86.67291162, 87.49012413, 80.80349802,
83.33959284, 83.33959284, 81.17452961, 84.6663701, 87.2275266,
82.33411513, 86.67291162, 81.17452961, 87.49012413, 86.67291162,
86.1228307, 87.2275266, 88.0609459, 87.2275266, 81.17452961,
82.33411513, 83.33959284, 82.33411513, 86.67291162, 81.17452961,
88.0609459])
```